

```
import pandas as pd

# Load datasets
customers = pd.read_csv('Customers.csv')
products = pd.read_csv('Products.csv')
transactions = pd.read_csv('Transactions.csv')

# Preview datasets
print(customers.head())
print(products.head())
print(transactions.head())
```

	CustomerID	CustomerName	Region	SignupDate
0	C0001	Lawrence Carroll	South America	2022-07-10
1	C0002	Elizabeth Lutz	Asia	2022-02-13
2	C0003	Michael Rivera	South America	2024-03-07
3	C0004	Kathleen Rodriguez	South America	2022-10-09
4	C0005	Laura Weber	Asia	2022-08-15

	ProductID	ProductName	Category	Price
0	P001	ActiveWear Biography	Books	169.30
1	P002	ActiveWear Smartwatch	Electronics	346.30
2	P003	ComfortLiving Biography	Books	44.12
3	P004	BookWorld Rug	Home Decor	95.69
4	P005	TechPro T-Shirt	Clothing	429.31

	TransactionID	CustomerID	ProductID	TransactionDate	Quantity \
0	T00001	C0199	P067	2024-08-25 12:38:23	1
1	T00112	C0146	P067	2024-05-27 22:23:54	1
2	T00166	C0127	P067	2024-04-25 07:38:55	1
3	T00272	C0087	P067	2024-03-26 22:55:37	2
4	T00363	C0070	P067	2024-03-21 15:10:10	3

	TotalValue	Price
0	300.68	300.68
1	300.68	300.68
2	300.68	300.68
3	601.36	300.68
4	902.04	300.68

```
# Merge transactions with products
transactions = transactions.merge(products, on='ProductID')

# Merge transactions with customers
data = transactions.merge(customers, on='CustomerID')

print(data.head())
```

	TransactionID	CustomerID	ProductID	TransactionDate	Quantity \
0	T00001	C0199	P067	2024-08-25 12:38:23	1
1	T00112	C0146	P067	2024-05-27 22:23:54	1
2	T00166	C0127	P067	2024-04-25 07:38:55	1
3	T00272	C0087	P067	2024-03-26 22:55:37	2
4	T00363	C0070	P067	2024-03-21 15:10:10	3

	TotalValue	Price_x	ProductName	Category	Price_y \
0	300.68	300.68	ComfortLiving Bluetooth Speaker	Electronics	300.68
1	300.68	300.68	ComfortLiving Bluetooth Speaker	Electronics	300.68
2	300.68	300.68	ComfortLiving Bluetooth Speaker	Electronics	300.68
3	601.36	300.68	ComfortLiving Bluetooth Speaker	Electronics	300.68
4	902.04	300.68	ComfortLiving Bluetooth Speaker	Electronics	300.68

	CustomerName	Region	SignupDate
0	Andrea Jenkins	Europe	2022-12-03
1	Brittany Harvey	Asia	2024-09-04
2	Kathryn Stevens	Europe	2024-04-04
3	Travis Campbell	South America	2024-04-11
4	Timothy Perez	Europe	2022-03-15

```
# Calculate total spending per customer
customer_spending = data.groupby('CustomerID')['TotalValue'].sum().reset_index()
customer_spending.rename(columns={'TotalValue': 'TotalSpending'}, inplace=True)
```

```
# Calculate average transaction value per customer
customer_avg_transaction = data.groupby('CustomerID')['TotalValue'].mean().reset_index()
customer_avg_transaction.rename(columns={'TotalValue': 'AvgTransactionValue'}, inplace=True)

# Count product category preferences
category_preferences = pd.crosstab(data['CustomerID'], data['Category'])

# Combine all features into a single DataFrame
customer_features = customers.merge(customer_spending, on='CustomerID', how='left')
customer_features = customer_features.merge(customer_avg_transaction, on='CustomerID', how='left')
customer_features = customer_features.merge(category_preferences, on='CustomerID', how='left')

# Fill missing values with 0
customer_features.fillna(0, inplace=True)

print(customer_features.head())
```

	CustomerID	CustomerName	Region	SignupDate	TotalSpending	\
0	C0001	Lawrence Carroll	South America	2022-07-10	3354.52	
1	C0002	Elizabeth Lutz	Asia	2022-02-13	1862.74	
2	C0003	Michael Rivera	South America	2024-03-07	2725.38	
3	C0004	Kathleen Rodriguez	South America	2022-10-09	5354.88	
4	C0005	Laura Weber	Asia	2022-08-15	2034.24	
	AvgTransactionValue	Books	Clothing	Electronics	Home Decor	
0	670.904	1.0	0.0	3.0	1.0	
1	465.685	0.0	2.0	0.0	2.0	
2	681.345	0.0	1.0	1.0	2.0	
3	669.360	3.0	0.0	2.0	3.0	
4	678.080	0.0	0.0	2.0	1.0	

```
from sklearn.preprocessing import MinMaxScaler

# Select numerical columns for normalization
numerical_cols = customer_features.select_dtypes(include=['float64', 'int64']).columns

# Normalize numerical columns
scaler = MinMaxScaler()
customer_features[numerical_cols] = scaler.fit_transform(customer_features[numerical_cols])

print(customer_features.head())
```

	CustomerID	CustomerName	Region	SignupDate	TotalSpending	\
0	C0001	Lawrence Carroll	South America	2022-07-10	0.314274	
1	C0002	Elizabeth Lutz	Asia	2022-02-13	0.174514	
2	C0003	Michael Rivera	South America	2024-03-07	0.255332	
3	C0004	Kathleen Rodriguez	South America	2022-10-09	0.501681	
4	C0005	Laura Weber	Asia	2022-08-15	0.190581	
	AvgTransactionValue	Books	Clothing	Electronics	Home Decor	
0	0.507057	0.2	0.0	0.6	0.166667	
1	0.351956	0.0	0.4	0.0	0.333333	
2	0.514948	0.0	0.2	0.2	0.333333	
3	0.505890	0.6	0.0	0.4	0.500000	
4	0.512480	0.0	0.0	0.4	0.166667	

```
from sklearn.metrics.pairwise import cosine_similarity

# Compute cosine similarity matrix
similarity_matrix = cosine_similarity(customer_features[numerical_cols])

# Convert to a DataFrame for easy manipulation
similarity_df = pd.DataFrame(similarity_matrix, index=customer_features['CustomerID'], columns=customer_features['CustomerID'])

print(similarity_df.head())
```

	CustomerID	C0001	C0002	C0003	C0004	C0005	C0006	\
	CustomerID							

C0001	1.000000	0.500277	0.808600	0.856999	0.951308	0.711859
C0002	0.500277	1.000000	0.885126	0.586321	0.591543	0.736503
C0003	0.808600	0.885126	1.000000	0.778289	0.889464	0.832574
C0004	0.856999	0.586321	0.778289	1.000000	0.758583	0.810212
C0005	0.951308	0.591543	0.889464	0.758583	1.000000	0.726567

CustomerID	C0007	C0008	C0009	C0010	...	C0191	C0192	\
CustomerID					...			
C0001	0.930154	0.802864	0.580197	0.395416	...	0.854094	0.928400	
C0002	0.612147	0.769940	0.767015	0.804178	...	0.356461	0.587797	
C0003	0.900881	0.849063	0.702609	0.623572	...	0.594605	0.783015	
C0004	0.750220	0.830519	0.411692	0.442937	...	0.878161	0.775528	
C0005	0.993777	0.754367	0.597472	0.387383	...	0.712469	0.868243	

CustomerID	C0193	C0194	C0195	C0196	C0197	C0198	\
CustomerID							
C0001	0.593943	0.898720	0.709999	0.714257	0.954553	0.800502	
C0002	0.374319	0.698138	0.949697	0.813518	0.585653	0.730442	
C0003	0.494673	0.846411	0.967996	0.918740	0.885190	0.853774	
C0004	0.787605	0.963363	0.744878	0.775095	0.759178	0.580685	
C0005	0.484430	0.810961	0.770387	0.789633	0.999660	0.863398	

CustomerID	C0199	C0200
CustomerID		
C0001	0.919925	0.788766
C0002	0.631991	0.844707
C0003	0.894424	0.921100
C0004	0.807205	0.776745
C0005	0.950526	0.811187

[5 rows x 200 columns]

```
# Function to get top 3 similar customers
def get_top_3_similar(customers_df, customer_id):
    similar_customers = customers_df[customer_id].sort_values(ascending=False).iloc[1:4]
    return list(zip(similar_customers.index, similar_customers.values))

# Generate recommendations for customers C0001 to C0020
lookalike_data = {}
for customer_id in customer_features['CustomerID'][:20]:
    lookalike_data[customer_id] = get_top_3_similar(similarity_df, customer_id)

# Convert to a DataFrame
lookalike_df = pd.DataFrame({
    'CustomerID': lookalike_data.keys(),
    'Lookalikes': [value for value in lookalike_data.values()]
})

print(lookalike_df)
```

	CustomerID	Lookalikes
0	C0001	[(C0146, 0.9849093056855885), (C0035, 0.984569...
1	C0002	[(C0134, 0.9861738653029506), (C0133, 0.976658...
2	C0003	[(C0166, 0.9981368903264345), (C0158, 0.990010...
3	C0004	[(C0017, 0.9721830996350582), (C0113, 0.969024...
4	C0005	[(C0197, 0.9996595915950468), (C0007, 0.993776...
5	C0006	[(C0135, 0.9910796141297742), (C0167, 0.983186...
6	C0007	[(C0005, 0.9937765536360452), (C0197, 0.990532...
7	C0008	[(C0162, 0.9656006917907872), (C0181, 0.945891...
8	C0009	[(C0034, 0.951112236091614), (C0092, 0.9490201...
9	C0010	[(C0077, 0.9895437462584057), (C0083, 0.979444...
10	C0011	[(C0126, 0.9921044763819368), (C0027, 0.989205...
11	C0012	[(C0065, 0.9876468679598578), (C0152, 0.983406...
12	C0013	[(C0107, 0.9896481906759469), (C0105, 0.988992...
13	C0014	[(C0151, 0.9852344025310596), (C0128, 0.968680...
14	C0015	[(C0123, 0.9969369053054044), (C0073, 0.972094...
15	C0016	[(C0183, 0.9999951812780966), (C0107, 0.999447...
16	C0017	[(C0075, 0.9883295070302321), (C0194, 0.986684...
17	C0018	[(C0168, 0.9788664595481981), (C0125, 0.973593...
18	C0019	[(C0191, 0.9571261980012316), (C0121, 0.942932...
19	C0020	[(C0130, 0.9909853390997839), (C0140, 0.970342...

```
lookalike_df.to_csv('Lookalike.csv', index=False)
```

Start coding or generate with AI.